

# Project Name: SwiftVisa AI-Based Visa Eligibility Screening Agent

Project summary (one line)

SwiftVisa is a retrieval-augmented screening agent that indexes official visa-policy documents per country into a vector database (FAISS) so the agent can quickly retrieve relevant policy chunks for eligibility-checking or question answering.

## Milestone 2: Build the RAG Pipeline and User Interfaces

### Goal

Implement the complete Retrieval-Augmented Generation (RAG) pipeline to enable grounded question answering, incorporate confidence scoring, and deploy user-friendly interfaces (CLI and Streamlit) for system interaction and evaluation.

### List of steps to achieve the goal:

1. **Retrieval Logic:** Implement the query embedding and FAISS similarity search.
2. **Prompt Engineering:** Design the system prompt and user context template for grounded generation.
3. **LLM Integration:** Integrate the Gemini/OpenAI client for structured JSON output.
4. **Confidence Scoring:** Develop a method to calculate and report a blended confidence score.
5. **Interfaces:** Build the CLI and Streamlit dashboard.
6. **Evaluation:** Implement batch query processing for testing and performance analysis.

### Implementation Steps:

1. **Retriever (rag/retriever.py):**
  - Accepts a user query.
  - Embeds the query using the same SentenceTransformer model used for indexing.
  - Performs a Nearest Neighbor (kNN) search against the persisted FAISS index to retrieve the top K relevant document chunk IDs and their similarity scores.
  - Maps the IDs back to the original chunk text and metadata using the persisted JSON file.
2. **Prompt Builder (rag/prompt\_builder.py):**
  - Constructs a comprehensive system instruction that enforces the **Analyst Persona** and demands **structured JSON output** for eligibility analysis.
  - Creates the final prompt by injecting the user's original query, and the retrieved

document chunks, framing the chunks as the "only source of truth."

3. **LLM Client (rag/llm\_client.py):**

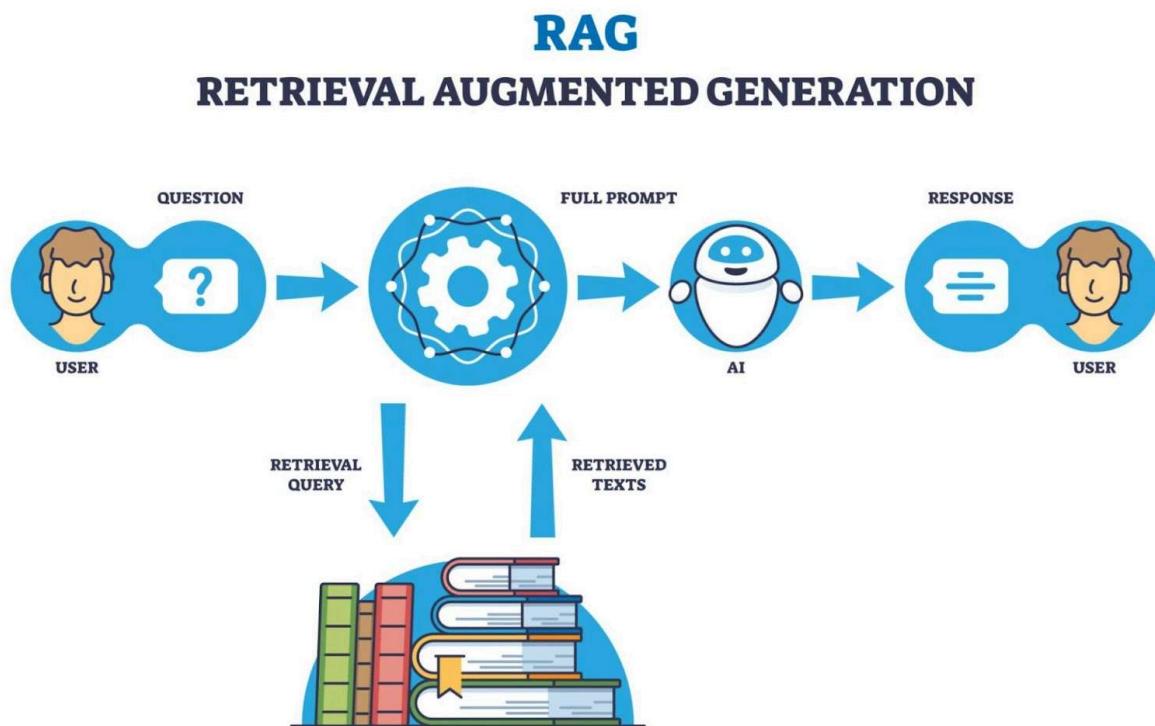
- Wraps the Gemini API calls.
- Enforces the output schema using the LLM's **Structured Output** feature (JSON schema).
- Implements a retry mechanism and robust try/catch logic to handle common API errors and invalid JSON responses.

4. **Pipeline (rag/pipeline.py):**

- Acts as the orchestrator, chaining the retrieval and generation components.
- Calculates the **Blended Confidence Score** (e.g., combining the max retrieval score with the LLM's self-reported confidence).
- Logs the query, retrieved chunks, and final answer using the rag/logger.py.

5. **User Interfaces:**

- **CLI (query\_cli.py):** Provides a simple, interactive console interface for immediate testing.
- **Dashboard (streamlit\_app.py):** Creates a responsive web interface for easier RAG testing and visual display of the structured response and retrieval metadata.



## Explore

### Deliverables produced (code files)

Files created or updated in this milestone:

File Name	Purpose and Function
rag/retriever.py	Implements the query embedding and FAISS search to get relevant chunks and scores.
rag/prompt_builder.py	Responsible for generating the strict system prompt and injecting RAG context.
rag/llm_client.py	Wrapper for LLM APIs, handles JSON schema enforcement and error resilience.
rag/pipeline.py	The main RAG execution logic; orchestrates retrieval, generation, and confidence scoring.
rag/logger.py	Captures queries, retrieved content, and final structured output for logging/auditing.
query_cli.py	Command-line interface allowing users to input profile data and queries.
streamlit_app.py	Streamlit Dashboard for an enhanced, visual user experience (Live Test Console).
process_test_queries.py	Script to run batch queries from user_queries.json and store structured results in query_results.json.

### Learnings from Milestone 2:

1. **Structured LLM Generation:**
  - o Understanding the importance of enforcing a JSON schema in the prompt/API call to ensure predictable, machine-readable output (decision, reason, confidence, future\_steps). This is critical for reliable downstream processing and display.
2. **RAG Efficacy and Prompting:**
  - o Mastering the RAG prompt template to instruct the LLM to prioritize the *provided context* (retrieved chunks) over its general knowledge, thereby guaranteeing a

grounded response.

### 3. Confidence Metrics:

- Implementing a combined confidence score—blending the quantitative FAISS retrieval score (how relevant the documents are) with the LLM's qualitative assessment (how confident it is in the final answer)—to provide a robust reliability metric to the user.

### 4. Deployment and Interface:

- Gaining experience in transitioning a core Python script into user-facing applications (CLI for developers, Streamlit for business users).

## Common pitfalls & how I handled them

Problem	Symptom	Fix implemented
<b>LLM Hallucination</b>	LLM ignores retrieved facts and relies on outdated internal knowledge.	Implemented a <b>strict system instruction</b> demanding the LLM reference the provided context <i>only</i> and to state "I cannot answer this based on the provided documents" if context is insufficient.
<b>JSON Parsing Failure</b>	LLM outputs text instead of the required JSON structure, breaking the pipeline.	Used the LLM API's native <b>Structured Output</b> feature to enforce the JSON schema. Added fallback Python try-catch and up to 3 retries with a corrective meta-prompt if parsing failed.
<b>Irrelevant Retrieval</b>	FAISS returns chunks with a high score that are semantically irrelevant to the <i>specific</i> question.	Reduced the maximum chunk size during indexing (Milestone 1) and fine-tuned the retrieval \$K\$ value in rag/retriever.py to balance precision and recall. Incorporated the blended confidence score to flag low-quality answers.

<b>API Rate Limits</b>	Scripts failed during batch testing due to exceeding API request limits.	Implemented a robust <b>exponential backoff</b> strategy in <code>rag/llm_client.py</code> to automatically wait and retry failed API calls, ensuring high availability during large batch runs.
------------------------	--	--